

Analysis of IMDb Movie Revenue

The Four-mula: Hellen Han, Lila Rogers, Amy Duan, Eva Aggarwal

2025-03-20

Introduction:

This project aims to examine key factors influencing IMDb movie revenues to understand what drives box office success. Revenue, the most widely used measure of success in the film industry, reflects audience demand and commercial viability. Prior research has highlighted various influences, such as star power, genre, and marketing, but the relative importance of these factors remains debated.

One central research question our project aims to answer is: **What production and release factors have the greatest impact on a movie's total revenue?** By addressing this question, we aim to deepen the understanding of revenue drivers and improve predictive models for box office performance.

We obtained our data set from [Kaggle](#), an online data science platform with a collection of community-developed open data sets. This data was collected by **Anand Shaw** from the **IMDb website** using various IMDb sites, and converted into a .csv file. The data was updated on a daily basis until 2 months ago.

Exploratory Data Analysis:

To get more indepth into our data set, [the data set](#) we used collects information available on the IMDb website for different movies, such that each observation describes characteristics of a specific movie. In general, the characters being measured follow basic information about the movie, various classifications of the movie's popularity and rating, and the monetary values associated with the movie. It measures the following **15 characteristics** per movie:

title: The name of the movie.

vote_average: The average rating the movie has received from users (on a scale, typically from 0 to 10).

vote_count: The total number of votes or ratings submitted for the movie.

status: The current state of the movie (e.g., “Released,” “Post-Production”).

release_date: The date when the movie was officially released.

revenue: The total earnings the movie made (usually in USD).

runtime: The duration of the movie in minutes.

adult: Indicates whether the movie is classified as adult content (e.g., “True” or “False”).

budget: The total cost of producing the movie (usually in USD).

imdb_id: The unique identifier for the movie on IMDb (Internet Movie Database).

original_language: The language in which the movie was originally produced (e.g., “en” for English).

popularity: A metric indicating how popular the movie is (typically based on views, searches, or ratings).

genres: The categories or genres the movie belongs to (e.g., Action, Comedy, Drama).

production_countries: The countries where the movie was produced.


spoken_languages: The languages spoken in the movie.

Due to the large size of the original data set and for the purpose of uploading this to our repository, we intentionally drop some characteristics, including `id`, `original_title`, `tagline`, `production_companies`, `overview`, `keyword` from the data set due to their redundancy, irrelevance to the research question and presence of a significant number of null values.

To obtain informative interpretations of the distribution of the revenue response variable, we analyze IMDb movies released from 2000 and beyond that have a revenue greater than 0 (non-missing values in this dataset). Furthermore, we log transform revenue, as the revenue generated across the IMDb movies is heavily right-skewed. (hellen could you include this in the description of the variables section, maybe write a short blurb for log transformation of log transformation for budget too).


Univariate EDA

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.000	9.933	15.323	13.387	17.646	21.822




The histogram shows the distribution of the log transformed IMDb movie revenues response variable, which is left-skewed and unimodal, indicating that on average, movies listed on IMDb tend to generate higher revenues. Furthermore, the distribution of the log transformed IMDb movie revenues has a center of approximately 1,800,000 US dollars, described by the median, and a spread of approximately 18,600,000 US dollars, described by the interquartile range. The majority of IMDb movies generate between approximately 24,100 and 18,600,000 US dollars.

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.000	9.629	15.425	13.222	17.034	20.500



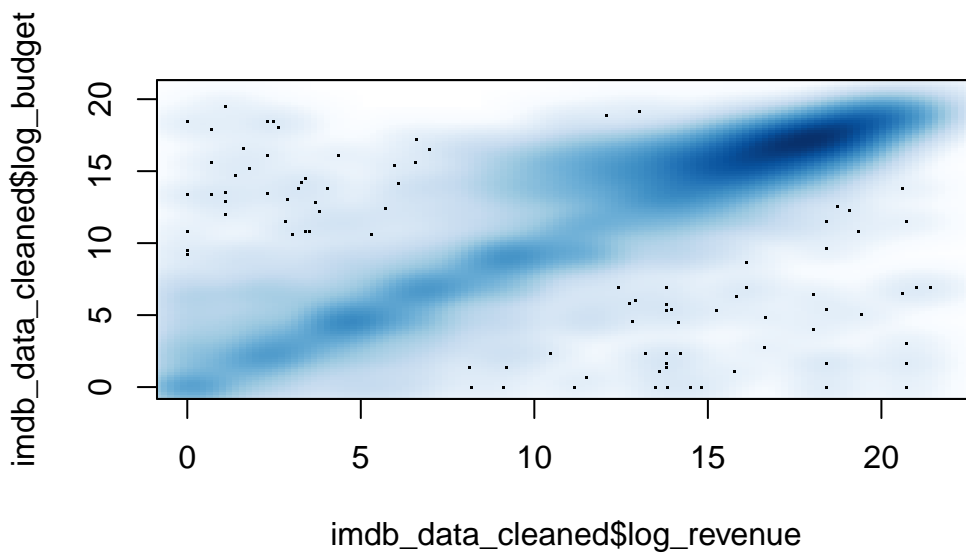

The histogram shows the distribution of the log transformed IMDb movie budget predictor variable, which is left-skewed and unimodal, indicating that although IMDb movie budget varies greatly, on average, movies listed on IMDb tend to have higher budgets. Furthermore, the distribution of the log transformed IMDb movie budget has a center of approximately 5,000,000 US dollars, described by the median, and a spread of approximately 25,000,000 US dollars, described by the interquartile range.

Distribution of IMDb Average Rating



The majority of IMDb movies' average ratings are concentrated between a rating of 4.0 to 8.0. A few IMDb movies have a low rating (average rating below 2.5), contributing to the left-skewness of the distribution of IMDb movies' average rating predictor variable. Furthermore, a few IMDb movies have an extremely high rating (average rating of 10.0).

Bivariate EDA




```
# A tibble: 2 x 5
  term      estimate std.error statistic  p.value
<chr>      <dbl>    <dbl>     <dbl>    <dbl>
```

1 (Intercept)	0.595	0.0649	9.16	6.08e-20
2 log_budget	0.967	0.00456	212.	0

[1] 0.8365862

From our visualization plotting the log of movie revenue versus the log of movie budget as well as our relatively high r-squared value associated with a linear regression model, we can identify that there is a strong linear correlation between log of movie revenue and log of movie budget. Our p-value is effectively 0, further exemplifying that there is a linear relationship. We expected this, as high budget movies are often more anticipated and therefore more people buy tickets. Because our data set is so large, we decided to also display this relationship using a smoother scatter visualization, which illustrates point density seen in our scatter plot.



In this visualization, we transformed the `vote_average` variable into a categorical variable taking levels “Low” for ratings 0-3, “Medium” for ratings 3-7, and “High” for ratings 7-10. From our boxplot, we can see that the median log revenue for Medium rated movies (around 16) is actually slightly higher than those with High ratings (15.5). The median log revenue for Low rated movies is significantly lower (6). The IQR of movies with high ratings is also much larger than for Medium and Low rated movies.

Potential Interaction Effects

Budget and Popularity —>

A tibble: 4 x 5

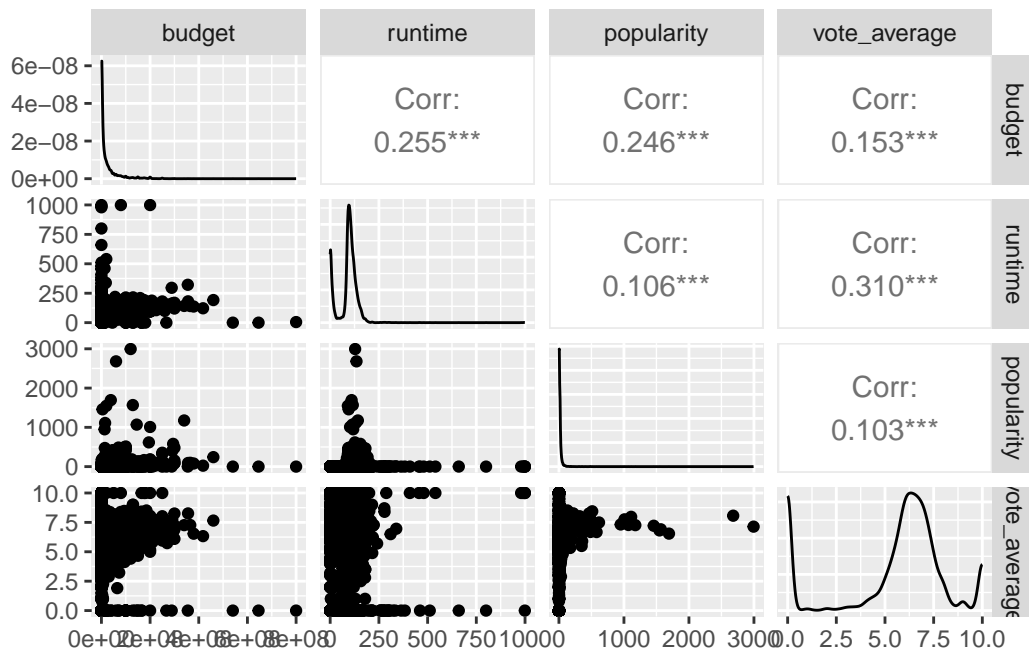
	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	-4.71e+6	1.42e+6	-3.33	0.000880
2	budget	2.98e+0	3.10e-2	95.9	0
3	popularity	8.44e+4	2.89e+4	2.92	0.00346
4	budget:popularity	1.71e-4	2.42e-4	0.706	0.480

A tibble: 4 x 5

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	10071338.	2855920.	3.53	4.23e- 4
2	vote_average	-1527974.	467222.	-3.27	1.08e- 3
3	vote_count	102245.	4249.	24.1	5.79e-124
4	vote_average:vote_count	-8150.	573.	-14.2	2.33e- 45

A tibble: 4 x 5

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	6962755.	1406895.	4.95	0.000000760
2	is_adult	NA	NA	NA	NA
3	vote_count	41859.	437.	95.9	0
4	is_adult:vote_count	NA	NA	NA	NA



! Important

Before you submit, make sure your code chunks are turned off with `echo: false` and there are no warnings or messages with `warning: false` and `message: false` in the YAML.